
The Floor: What Specification Cannot Remove

Wenhao Lu
NXT Horizon
wenhao@nxthorizon.org

Abstract

The Floor is a three-minute video work built from 480 generations: 240 images from a diffusion model and 240 texts from a language model, produced under identical, fully controlled conditions. In three movements it stages a single question as a visual argument: when an artist specifies a work in ever greater detail, what part of the result was never theirs to specify? The piece watches one prompt grow from three words to thirty-five while the images refuse to converge; it holds one random seed fixed across twelve unrelated prompts until the seed’s own faint handwriting becomes visible in where the light falls; and it closes on an asymmetry between the two media, in which the image’s randomness can be summoned again by name while the sentence’s cannot. The floor of the title is the share of every generated artifact that specification cannot remove, authored by no one, and always present in the room where the work is made.

1 The work

The Floor is a single-channel video, 2 minutes 58 seconds, 1920×1080, silent text on black between full-frame image sequences. Every visual element in the piece is a generation or a direct arrangement of generations; nothing is illustrated, simulated, or staged. The material is a fully balanced grid: two prompt families (a teapot, a coastal landscape), six nested prompt “rungs” per family growing from 3 to 35 words, ten fixed seeds, two diffusion models, plus a paired set of 240 text generations from a language model under the same twelve prompts. The piece unfolds in three movements.

I. The floor. One seed is held fixed while the prompt grows rung by rung, each new constraint highlighted as it arrives: *a teapot* → *a ceramic teapot* → *a blue ceramic teapot on a table* → ... → a 35-word product-photograph specification excluding shadows, other objects, and text. The images narrow. They do not close. At the movement’s end a card states the measurement the sequence has just performed: *thirty-four words removed sixty-nine per cent of it; the rest is not the artist’s, and never was*. The escalating prompt is the familiar ritual of control; the residual spread is the piece’s subject.

II. The hand. The grid is turned ninety degrees: now one seed is held fixed across all twelve unrelated prompts, teapots above, coastlines below (Figure 1). Between rows, the piece shows each seed’s mean luminance map, its average distribution of light over every prompt, as a small grey glyph. Watched in sequence, the glyphs differ from seed to seed and persist within one: *the seed does not choose what you get; it chooses where the light falls, and it chooses the same way every time*. The randomness, in other words, has a hand, faint, stable, and unchosen, and the movement makes that hand visible without a single number on screen.

III. Addressable. The final movement pairs the two media at their most constrained: the same maximal prompt sent twenty times to a diffusion model with the seed pinned, and twenty times to a language model with the seed pinned. On the left, the image returns identical every time, *0 pixels differ from the last call*. On the right, each sentence arrives roughly forty per cent new, because the serving infrastructure silently discards the seed parameter (Figure 2). The closing card: *Same*




Figure 1: Still from movement II, *the hand*: twelve unrelated prompts, one seed. The seed does not decide the subject, but it leans on palette, geometry, and the placement of light in a way that recurs across every prompt.




Figure 2: Still from movement III, *addressable*: the same fixed prompt and fixed seed, called repeatedly. The diffusion model returns the identical image; the language model's provider discards the seed, so each call produces a new text.

randomness. One of them you can hold. Every image in the piece can be made again; none of the sentences can.

2 The roles of AI and ML

Generative models are both the medium and the subject. All 240 images were produced by two SDXL-class diffusion models (Stable Diffusion XL 1.0, Playground v2) via API with every sampler setting held constant, so that within each cell the seed is the only free variable; all 240 texts by an open-weight language model (gpt-oss-120b) at temperature 1.0 under the same twelve prompts. The piece's editorial claims are not rhetorical: they come from a variance decomposition run over the same 480 artifacts, in which within-cell spread is, by construction, the randomness the prompt does not control. The "sixty-nine per cent" of movement I, the per-seed luminance glyphs of movement II, and the "forty per cent of the words are new" counter of movement III are each computed from the

material on screen. ML also appears once behind the scenes: the semantic analysis that produced the quoted figure embeds machine captions of the images, a model reading a model, which the artist then reads in turn. The work is thus a case of measurement used as composition, with the statistics supplying the piece's dramaturgy rather than a supplementary appendix.

3 How the theme of Agency is addressed

The piece proposes that agency in generative practice is not a spectrum from tool to collaborator but a *partition*, and it makes the partition visible.

Movement I addresses the artist's agency at its most earnest: the prompt as an instrument of will, sharpened word by word. The movement grants that instrument everything, thirty-five words, exclusion clauses, compositional directives, and then shows the remainder that no clause reached. That remainder is the floor: a share of every generated work that is authored by no one, since the artist cannot intend it, the model's makers did not choose it, and the system merely resolves it.

Movement II complicates the remainder. If the leftover randomness were mere noise, the artist's relation to it would be a lottery. But the seed's faint hand, its stable way of placing light across unrelated subjects, means the everyday practice of regenerating until something clicks is not sampling but *casting*: a selection among anonymous collaborators with dispositions, none of which the artist wrote and one of which the artist can choose. Agency reappears here in a diminished but genuine form: not the power to intend the remainder, but the power to recognise and elect it.

Movement III is about who grants even that. Whether randomness is addressable at all, whether a seed can be revisited, a result reproduced, a collaborator recalled by name, turns out to be a property of serving infrastructure, decided by a platform and changeable without notice. On one side of the frame the artist retains a handle on chance; on the other, the same gesture of fixing a seed is silently ignored, and the stochastic collaborator becomes unnameable as well as unchosen. The piece ends on this asymmetry because it is where the track's question lands hardest: some of the agency being negotiated in creative AI is not between artist and model at all, but between both of them and the infrastructure in the middle.

4 Author biography

Wenhao Lu is a researcher and poet working at the intersection of AI and the humanities, interested in both directions of that exchange: how AI can revive and reanimate humanistic practice, and how the close-reading habits of the humanities can inspire new approaches to AI interpretability. He is a National Student Poet Program semifinalist, selected among the top 50 of roughly 36,000 Scholastic submissions, and his poetic practice informs the questions this work asks about agency. He works with NXT Horizon, an AI incubator for underrepresented students.