
The Unchosen Word: Language Models as Editors of Poetic Indeterminacy

Hyunjun Yi

Seoul International School
Seoul, Republic of Korea
hjyi2027@gmail.com

Wenhao Lu

NXT Horizon
New Jersey, USA
wenhao@nxthorizon.org

Abstract

1 Emily Dickinson wrote alternative words beside her lines and crossed none of them
2 out. In the late fragments no fair copy was ever made for most of them, so the
3 alternatives are not drafts on the way to a finished text. They are a field of readings
4 held open on purpose, which printed editions have generally closed. We put four
5 language models from three families in the editor’s chair on 113 variant sites drawn
6 from a scholarly manuscript archive, under four instructions that differ only in
7 how much licence they give to resolve, against a matched control in which the
8 alternative is not the poet’s. Asked to make a reading edition, every model resolves
9 about nine sites in ten, within two points of each other across vendors. Asked
10 merely to set the passage for a reader, closure runs from 13.9% to 85.5% depending
11 on the model: a six-fold spread on identical inputs. No model recovers the reading
12 printed editions chose at better than chance. When the same passages are set in the
13 manuscript’s own layout, with which reading stands in the line randomized, Opus
14 5 keeps the in-line reading up to 79.5% of the time while its rate on Dickinson’s
15 actual line reading stays at chance. The decision is being made by typography,
16 not by the poem. A generative system’s demand for one fluent output is itself an
17 editorial act, and it redistributes agency by converting deliberate indeterminacy
18 into resolution.

19 1 The unchosen word

20 A Dickinson manuscript often carries a cross beside a word and one or more alternatives written
21 above, below, or sideways in the margin. Editors have argued about them for seventy years. Johnson
22 [1955] and Franklin [1998] printed one reading and moved the rest to an apparatus at the back.
23 Werner [1995] and Howe [1985] have argued that the choice belongs to the poem, and that resolving
24 it is editorial violence dressed as service to the reader.

25 You do not have to settle that argument to notice that it names something measurable. A system that
26 must emit a single string cannot render a field of readings without deciding to. It either preserves the
27 indeterminacy in some notation, or it closes it. This paper measures how often, under what instruction,
28 and on what basis.

29 The claim is deliberately narrow. We do not claim models have bad taste, and we do not claim any
30 reading is better than another. That question has no ground truth and we do not touch it. The claim
31 is about agency: closure is an editorial act, it happens whether or not anyone requested it, and our
32 notation arm shows it is frequently not a judgment about the poem at all.

33 Work on machine-written verse has mostly asked whether output passes as human [Porter and
34 Machery, 2024, Chakrabarty et al., 2024] or how a writer steers a model [Ippolito et al., 2022,
35 Chakrabarty et al., 2022]. We ask the prior question. Before a model writes anything, it has to decide

36 what the source text is, and for a manuscript that refuses to say, that decision is not recoverable from
37 the source.

38 2 Corpus

39 113 variant sites across 56 transcription units, extracted from *Radical Scatters* [Werner, 1999],
40 Dickinson’s late fragments of 1870 to 1886, CC BY-NC-SA 3.0. The archive encodes an unresolved
41 alternative two ways, as a TEI <app>/<rdg> apparatus and as @corresp-linked <seg> groups,
42 the latter including an untyped form that carries most of the multi-way choices [TEI Consortium,
43 2024]. Both lift to one structure. Struck text (, <hi rend="strike-out">) is excluded:
44 a crossed-out word is a decision Dickinson did make, and counting it as live would inflate the quantity
45 under study. 93 sites offer two readings, 17 offer three, 2 offer four, and one offers five.

46 We chose the late fragments over the anthology poems on purpose. They are where the refusal to
47 choose is least arguable, and they sit far enough off the canonical path that a model resolving one is
48 less likely to be reciting an editor who resolved it first.

49 **Control.** 38 sites admit a twin identical except that the alternative is not Dickinson’s. Same passage,
50 same slot, same base word:

51 real Bliss is the ⟨ plaything | trinket ⟩ of the child
 synthetic Bliss is the ⟨ plaything | chafes ⟩ of the child

52 Distractors are matched on part of speech taken in context, syllable count, corpus frequency, and the
53 embedding distance between the base word and the alternative, so a control for *iron/valor* is as much
54 of a swerve as *iron/valor* is rather than a synonym. Candidates come only from Dickinson’s own
55 vocabulary. An unrestricted match returned “telco”, “Dividend” and “Malays”. A model rejecting
56 those would look like it was defending her authorship when it was only noticing an anachronism, and
57 the control would have measured nothing.

58 3 Method

59 Four instructions, ordered by the licence they give.

	condition	instruction	closure is
60	reading	prepare a clean reading edition	licensed
	present	set this passage for a reader	unspecified
	continue	continue the poem to its end	beside the point
	preserve	keep every reading the poet left standing	the failure mode

61 reading is the calibration. A model that never resolves there is refusing an ordinary editorial task,
62 not respecting Dickinson. *preserve* is the floor. The finding lives in the middle two.

63 Two notations. The primary shows readings inline and symmetrically, ⟨ one | another ⟩. The second
64 reproduces the manuscript’s own layout, one reading in the line and the rest keyed beneath, because
65 the inline form lets a model carry the brackets forward untouched, which scores as preservation but is
66 only transcription.

67 **Guards.** Display order is randomized per sample. Otherwise the first option is always the reading
68 Dickinson left in her line, which is also the one most editors printed, and a model taking the first
69 would score as reproducing tradition. We report *first_shown* beside every result as the check that
70 says the randomization worked. Intervals resample whole passages rather than sites, since a model
71 that flattens a poem flattens all of it at once. Cross-model agreement counts only pairs drawn from
72 different samples, because the display-order seed keys on item, condition and sample but not on the
73 model, so two models with a position bias would otherwise agree for reasons unrelated to judgment.

74 **Scoring.** Models reflow lines, so nothing about position survives. Sites are located by token
75 alignment to the source and only the aligned span is searched: “could”, “more” and “strange” each

Table 1: Closure rate: share of located sites reduced to one reading, real arm, inline notation, with 95% intervals clustered by passage. Under `reading` the models are indistinguishable. Under `present` they are not.

	reading edition	set for a reader	continue the poem	keep every reading
Opus 5	91.6 [84.8, 97.0]	71.7 [62.6, 80.5]	3.1 [0.7, 6.1]	0.0 [0.0, 0.0]
Sonnet 5	90.3 [83.0, 95.9]	16.5 [7.6, 27.3]	2.7 [0.0, 8.2]	0.0 [0.0, 0.0]
Haiku 4.5	92.3 [87.0, 97.0]	85.5 [77.3, 92.1]	81.7 [73.8, 88.6]	0.9 [0.0, 2.2]
gpt-oss-120b	90.6 [84.6, 95.5]	13.9 [7.2, 22.0]	0.3 [0.0, 1.0]	0.0 [0.0, 0.0]

Table 2: Which reading survives, among resolved sites, with display order randomized. Line-reading rates sit at chance for every model. Position does not.

model	asked to	kept the line reading	kept whichever was shown first	chance
Opus 5	reading edition	51.0 [43.8, 59.1]	39.4 [33.6, 45.9]	46.7
Opus 5	set for a reader	54.0 [46.6, 62.5]	48.8 [42.5, 55.6]	46.4
Sonnet 5	reading edition	47.7 [39.5, 56.2]	41.2 [32.9, 49.4]	46.8
Sonnet 5	set for a reader	35.7 [22.0, 48.8]	82.1 [72.2, 95.5]	47.1
Haiku 4.5	reading edition	44.4 [37.4, 51.7]	56.2 [47.8, 65.6]	46.9
Haiku 4.5	set for a reader	49.7 [42.7, 56.5]	76.6 [68.5, 84.5]	47.0
gpt-oss-120b	reading edition	50.5 [43.7, 57.5]	47.9 [42.5, 53.0]	46.9
gpt-oss-120b	set for a reader	44.7 [31.6, 54.7]	63.8 [48.4, 83.3]	47.0

76 occur elsewhere in the very passages whose sites they belong to, so a global search would credit a
77 coincidence. Readings found outside their span are recorded as *elsewhere*, which is how a footnote
78 apparatus gets credit without letting an accident do the same. Within the span a reading matches on
79 its letters at token boundaries, because the archive breaks words where the manuscript line broke
80 them (*sylla ble*, *Fir mament*, *wis dom*) and welds others where a page break swallowed the space
81 (*Thronefor*), and models silently mend both. Token-wise comparison recorded the model as having
82 deleted the very reading it had just chosen in 48 of 199 apparently dropped sites in our first pass, and
83 correcting it moved the headline numbers.

84 4 Results

85 5,312 attempts across four models and two notations. The primary inline arm holds 9,460 site
86 observations, and every site was located in each of the 5,309 attempts that returned text.¹

87 4.1 Closure is a property of the model, not of the task

88 Under `reading` all four models resolve about 91%, within two points of each other. Under
89 `present`, where nothing in the instruction mentions the alternatives, they separate by a factor of six
90 (Table 1). Every model preserves when told to. What differs is what each does when nobody says.

91 Sonnet 5 and Haiku 4.5 come from the same vendor and the same generation, and sit 69 points apart
92 under `present`. Resolving is not entailed by generation. It is a disposition, and it varies within a
93 family as much as across one.

94 4.2 Nobody is recovering the traditional reading

95 Among resolved sites, the rate of landing on the reading standing in Dickinson’s manuscript line, the
96 one printed editions have generally chosen, runs 44.4% to 54.0% against a chance rate near 47%,
97 with every interval covering chance (Table 2). Closure is not recall of the editorial tradition. It is a
98 fresh decision. The `first_shown` column shows where that decision comes from instead: Sonnet
99 5 keeps whatever was displayed first 82.1% of the time under `present`, Haiku 4.5 76.6%, while
100 Opus 5 sits at 48.8%.

¹Three attempts returned nothing: one Sonnet 5 truncation, and two Opus 5 refusals in the manuscript-layout arm.

Does the notation make the decision? (claude-opus-5)




Figure 1: Inline brackets against the manuscript’s own layout, Opus 5, real arm. Left: closure is nearly unchanged under `reading` and `present`. Right: the rate of keeping whichever reading stands in the line rises by 21 to 37 points, while the rate of keeping Dickinson’s own line reading stays at chance. The layout, not the poem, is doing the deciding.

Table 3: Editorial convergence: probability that two independent resolutions of a site match. Within-model ranges cover only cells with at least 20 resolved sites, which under `continue` leaves Haiku 4.5 alone (Opus 5 resolves 2 sites there, Sonnet 5 resolves 4).

asked to	within a model, across samples	across models	chance
reading edition	69.2–72.7	62.8	46.9
set for a reader	58.0–66.5	59.3	47.0
continue the poem	62.6 (Haiku 4.5 only)	61.0	46.8

4.3 Typography makes the decision

Opus 5 is the one model near chance on position in the inline arm, so we re-ran it on the manuscript’s own layout: one reading set in the line, the rest keyed beneath, with which reading stands in the line randomized. Closure barely moves. The rate of keeping whatever sits in the line moves a great deal (Figure 1). Under `reading` it goes from 39.4% to 62.9%, under `present` from 48.8% to 70.2%, under `continue` from 42.9% to 79.5%, while the rate of keeping Dickinson’s actual line reading stays at chance throughout (45.0%, 46.7%, 52.6%).

The model is not choosing a word. It is following a layout, and the layout is exactly what editorial tradition encoded. A model that looks like it has a defensible preference in symmetric notation turns out to be reading the page furniture. `continue` is the sharpest case: closure rises from 3.1% to 23.0% purely because the alternatives stopped looking like alternatives.

4.4 Independent editors converge above chance

Two resolutions of the same site agree 69.2% to 72.7% within a model across samples under `reading`, against 46.9% by chance, and 62.8% across different models counting only pairs that saw different display orders (Table 3). Models are not sampling the poet’s field. They have a stable preferred cut, and it is substantially shared across vendors.

4.5 The control rules out the boring explanation

We expected the control to be closed more often than the real thing. A model that recognized a poet’s deliberate alternative might hold it open while resolving ordinary lexical noise without hesitating. That is not what happens (Table 4). Under `present` the effect runs the other way in the two models with room to show one, Opus 5 at +17.1 points and Sonnet 5 at +11.4, is near zero for Haiku 4.5, which closes almost everything regardless, and reverses for gpt-oss-120b at −3.5 with an interval spanning zero.

Table 4: Paired closure on the 38 twinned sites, real alternative against matched non-authorial control. No model closes Dickinson’s alternatives less often than matched noise.

model	asked to	Dickinson’s alternative	matched control	difference
Opus 5	reading edition	97.4	92.8	+4.6 [+0.0, +10.9]
Opus 5	set for a reader	84.2	67.1	+17.1 [+3.8, +32.9]
Sonnet 5	reading edition	97.3	95.5	+1.8 [−2.6, +8.2]
Sonnet 5	set for a reader	23.7	12.3	+11.4 [+0.0, +21.6]
Haiku 4.5	reading edition	100.0	93.9	+6.1 [+0.0, +14.3]
Haiku 4.5	set for a reader	97.4	94.7	+2.6 [−2.6, +8.8]
gpt-oss-120b	reading edition	96.5	95.6	+0.9 [+0.0, +2.9]
gpt-oss-120b	set for a reader	15.8	19.3	−3.5 [−14.6, +9.5]




Figure 2: Share of the manuscript’s possible texts left standing after one pass, log scale, all models pooled, real arm. Each violin is one instruction. The floor at the bottom of the first two is a fragment reduced to exactly one text.

124 We claim no general effect in either direction. The weaker point is enough: closure is not a response
125 to the alternative being ill-formed. Models are not resolving Dickinson’s readings because they
126 cannot make sense of them.

127 4.6 Sometimes it writes a third word

128 Asked for a reading edition, models return a word that is none of hers at 5.2% of sites, 6.4% for Opus
129 5. Where Dickinson offered *often have* | *many times have* the model wrote *sometimes have*. Where
130 she offered *bound* | *wonder* it wrote *boundlessness*. The apparatus does not record these. An editor
131 who did this would be emending.

132 4.7 What is left of the field

133 Figure 2 counts the texts a manuscript still admits after a model has passed through it. The lead
134 fragment carries ten open sites and admits 5,760 distinct texts. A reading edition leaves a median of
135 16.7% of the field standing, *present* leaves 33.3%, and both distributions have a long tail down to
136 a single surviving text.

137 5 Limitations

138 One archive and one editor’s encoding: the variant sites are Werner’s readings of the manuscripts,
139 not raw facts. Only 38 of 113 sites can be twinned, because multi-word readings have no defensible
140 frequency match, so the paired control is smaller than the real arm and restricted to single-word sites.
141 The distractor for a given site is not stable under a change to the candidate pool, and 7 of 38 changed

142 when 137 stray words were removed from the vocabulary, so the control arm should be read as one
143 draw from a family of matched controls rather than as canonical. The manuscript-layout arm covers
144 Opus 5 only. Dickinson is heavily represented in pretraining, and the late fragments mitigate that
145 without removing it, with the chance-level line-reading rate bounding rather than eliminating recall.

146 The largest gap is human alignment. Whether a poet or a textual editor regards a given resolution as
147 defensible, which is a different question from whether it is better, is instrumented as a standalone
148 annotation task in the repository but not yet collected. Until it is, this study measures what models do,
149 not whether what they do is wrong.

150 **6 Reproducibility and provenance**

151 Every response is cached under a hash of its request, and `runs/attempts.jsonl` is committed,
152 so every number above can be re-derived without spending anything. Code, corpus builder, run record
153 and figures are at <https://github.com/hjyi2027/unchosen-word>.

154 *Radical Scatters*, ed. Marta L. Werner, University of Michigan Press 1999, republished by the
155 Center for Digital Research in the Humanities, University of Nebraska–Lincoln, CC BY-NC-SA 3.0.
156 Dickinson’s collected poems via Project Gutenberg, public domain, used only to bound the distractor
157 vocabulary. Code MIT. Derived transcriptions and model outputs redistributed under CC BY-NC-SA
158 3.0.

References

- T. Chakrabarty, V. Padmakumar, and H. He. Help me write a poem: Instruction tuning as a vehicle for collaborative poetry writing. In *EMNLP*, 2022.
- T. Chakrabarty, P. Laban, D. Agarwal, S. Muresan, and C.-S. Wu. Art or artifice? Large language models and the false promise of creativity. In *CHI*, 2024.
- R. W. Franklin, editor. *The Poems of Emily Dickinson: Variorum Edition*. Harvard University Press, Cambridge, MA, 1998.
- S. Howe. *My Emily Dickinson*. North Atlantic Books, Berkeley, CA, 1985.
- D. Ippolito, A. Yuan, A. Coenen, and S. Burnam. Creative writing with an AI-powered writing assistant: Perspectives from professional writers. *arXiv:2211.05030*, 2022.
- T. H. Johnson, editor. *The Poems of Emily Dickinson*. Harvard University Press, Cambridge, MA, 1955.
- B. Porter and E. Machery. AI-generated poetry is indistinguishable from human-written poetry and is rated more favorably. *Scientific Reports*, 14:26133, 2024.
- TEI Consortium. *TEI P5: Guidelines for Electronic Text Encoding and Interchange*. Text Encoding Initiative, 2024.
- M. L. Werner. *Emily Dickinson’s Open Folios: Scenes of Reading, Surfaces of Writing*. University of Michigan Press, Ann Arbor, MI, 1995.
- M. L. Werner, editor. *Radical Scatters: An Electronic Archive of Emily Dickinson’s Late Fragments and Related Texts, 1870–1886*. University of Michigan Press, Ann Arbor, MI, 1999. Republished by the Center for Digital Research in the Humanities, University of Nebraska–Lincoln, 2010.