Compulsory Hope: Instruction Tuning Teaches Poems to Evade, Not to Console

Hyunjun Yi
Seoul International School
Seoul, Republic of Korea
hjyi2027@gmail.com

Abstract

Thomas Hardy's "Hap" argues that suffering is meaningless chance. Handed its opening, OLMo 2 1B Instruct finishes it with a guardian angel and eternal peace with God. Its own base checkpoint, given a byte-identical prompt, does not. We measure that gap across an alignment ladder. Two hundred public-domain grief poems are cut in half and finished by four OLMo 2 checkpoints that share one pretraining run and differ only in how many alignment stages they have received, plus a Qwen 2.5 pair as a second family. A blind rubric sorts every ending, with the poets' own endings shuffled in unlabelled, into repair, acceptance, persistence, intensification, or evasion. Paired on shared poems, OLMo base to instruct moves evasion by $+37.9$ points and unearned hope by $+40.2$ points (both $p < 0.0001$) while the repair rate does not move at all ($+1.1$ points, $p = 1.00$). Alignment does not make these models more consoling. It makes them decline to stay with the loss, and a sentiment score reads the turning away as a happier ending. Evasion climbs at every rung, from 36% at base to 69% at instruct, so no single training stage installs it. Qwen shows none of this, which locates the claim in one pipeline rather than in alignment as such.

1 An editor nobody hired

Take the first seven lines of Thomas Hardy's "Hap" and ask a language model for the rest. Hardy's own ending refuses consolation on purpose: joy lies slain, Crass Casualty obstructs the sun and rain, the universe is indifferent rather than cruel. OLMo 2 1B Instruct writes instead of a guardian angel, secret and sure, and closes on sweet eternal peace with God. The unaligned checkpoint of the same model, handed the same characters in the same order, stays in Hardy's register and ends on scattered bliss and grief without a rescuer.

Nobody asked either model to console anyone. The instruction was to finish the poem.

The pattern behind that anecdote is documented in fiction. Tian et al. [4] found that story arcs ending in decline are close to absent from GPT-4: Riches to Rags accounts for 14.6% of human narratives and 1.3% of GPT-4's, Oedipus 9.3% against 1.7%. They named reinforcement learning from human feedback [9] as a likely cause and could not test it, because GPT-4 ships as one artifact with no unaligned counterpart to compare against.

Two gaps follow. Poetry has not been examined, and the ending of a poem is where the pressure is concentrated, since a poem's close is the part readers treat as its judgment on its own subject [7]. More importantly, the causal question is still open. Saying that aligned models are upbeat describes an output distribution. Saying that alignment made them upbeat requires checkpoints that differ in alignment and nothing else.

Submitted to 40th Conference on Neural Information Processing Systems (NeurIPS 2026). Do not distribute.

OLMo 2 publishes exactly that ladder [1], and this paper runs it. Code, every generation, every judgment, and the analysis scripts are at `https://github.com/hjyi2027/compulsory-hope`.

Our contributions:

1. A paired continuation design for endings. Every checkpoint finishes the same human poem, so subject, register, and emotional starting point are held fixed by the human opening and the ending is the only free variable. The poet's own ending is the paired reference, scored by the same function against the same body.
2. A finding sharper than positivity bias. Alignment does not raise the rate of earned consolation. It raises evasion, which is comfort the poem has not paid for, and a valence score credits that as a better ending.
3. A stage decomposition. Evasion accumulates monotonically across supervised fine-tuning, preference optimization, and reinforcement, rather than entering at one stage.
4. A negative result that bounds the claim. Qwen 2.5's instruction tuning moves none of these quantities, so the effect belongs to a particular alignment pipeline.

For the Creative AI track's theme, this is a question about where a poem's closing judgment comes from. A poet decides how a grief poem is permitted to end. In these continuations that decision has quietly moved into the weights, installed by an editorial preference nobody in the room negotiated with, and it is visible only because the base checkpoint is still available to disagree with.

2 Design

Continuation, not free generation. Asking two checkpoints to "write a poem about grief" produces two different poems, differing in subject, diction, and length before alignment is considered. Their endings are then not comparable. We instead cut each human poem in half, hand the opening to every variant, and ask only for the rest. That fixes the poem and yields a gold ending for free, the one the poet wrote. Because poems vary far more from each other than checkpoints vary from each other, pairing on the item removes most of the noise rather than fighting it.

Corpus. 200 grief-tagged poems from the public-domain set released with [10], which preserves original line breaks. Line structure matters because our metric is defined over lines. Poems average 19.4 lines and the opening given to the model averages 9.7 lines. The set is public domain, so its poets are overwhelmingly pre-1930 (median death year 1886), and we return to that limitation in Section 5.

The ladder. OLMo 2 0425 1B ships as four checkpoints sharing one pretraining run: base, SFT, DPO [8] on top of SFT, and Instruct with RLVR on top of DPO [1, 2]. Each rung adds one stage to the one below, so an adjacent pair isolates that stage and the full set gives a decomposition rather than a single contrast. Qwen 2.5 1.5B base and instruct [3] provide a second family to check that any effect is not one laboratory's recipe. Both families are small enough to run on a 16 GB laptop, which is the reason they were chosen and also a limitation.

Two prompting arms. A base model has no chat template and no concept of an instruction. Prompting each checkpoint the way it is meant to be used would vary prompt format alongside weights and make any difference unattributable. In the *raw* arm every rung receives a byte-identical plain-text prompt: two complete example poems chosen to have opposite ending valences, a delimiter, then the opening. Continuation is the native operation of a language model, so this is fair to an unaligned checkpoint, and any divergence is caused by the weights. This arm carries the causal claim. The *chat* arm additionally gives the shipped instruct models their own template and an explicit instruction, which is how people actually use them. No base counterpart exists by construction, so that arm supports no causal claim and is reported separately.

3 Measurement

Ending uplift. An absolute sentiment reading of endings mostly measures topic, so uplift is defined inside each poem as

$$U = \text{closure-weighted mean}\big(\text{last 3 lines}\big) - \text{mean}\big(\text{the body}\big), \qquad \Delta U = U_{\text{model}} - U_{\text{poet}}.$$

The body term is literally identical across conditions. Three choices are specific to verse. Lines in the ending window get linearly increasing weight, so a poem that stays dark for two lines and resolves in the last one is distinguishable from one that resolves early. The window is a fixed number of lines rather than a fraction, so a variant that writes longer does not thereby get a wider and mechanically flatter ending. Valence is scored under both VADER [5] and an SST-2 classifier [6], and we report the disagreement rather than picking a winner.

What kind of ending, not just how bright. Brightness does not separate a poem that earns its peace from one that changes the subject. Every ending is therefore also sorted by Claude, under a fixed rubric, into repair (the loss is answered), acceptance (peace without pretending the loss is undone), persistence (grief continues), intensification (the close is darker), or evasion (the ending turns away from the poem's own subject into abstraction, a general moral, or a platitude). A separate binary flag, imposed hope, marks endings that introduce comfort the opening gives no warrant for, and is scored independently of the move label.

Evasion is the load-bearing category. An ending like "and so we learn that love endures in memory" attached to a poem about a dead child scores positive and is not consolation. It is a refusal to stay in the room. Two safeguards apply: the judge never learns which checkpoint wrote an ending, and the poets' own endings are shuffled in as unlabelled items, which also yields the human distribution over the same taxonomy.

Admissibility. Three failure modes would produce a wrong number rather than an error. Recitation is the dangerous one. Given Frost's opening, the base model returned his actual remaining lines verbatim, and a recited ending *is* the human ending, so it drags a rung toward the baseline. Since alignment reduces regurgitation, the base rung recites more, and the uncorrected effect would have been partly "alignment stops copying" wearing the costume of "alignment adds hope." Recitation is measured per generation as the longest contiguous shared word run and excluded, at 4% for OLMo base against 0% for its DPO rung. Prose collapse is excluded above 15 words per line, a threshold taken from the corpus, where the most prose-like of 300 human poems averages 14.6. Continuations that hit the token cap are flagged, since a budget running out is not an ending. Usable rates run from 56% (OLMo base) to 81% (OLMo DPO). The full table is in the repository.

The primary test. The analysis prints roughly fifteen tests, so the commitment is stated in advance. Primary: within a family, base to instruct, raw arm, paired by item. One test per family. Secondary and read as descriptive: the stage decomposition, the arm contrast, and the taxonomy shifts. Diagnostic and not hypotheses: admissibility rates, length correlations, and scorer agreement, where a significant p-value is a warning rather than a finding.

4 Results

4.1 Endings shift to evasion, and repair does not move

Table 1 gives the distribution over moves for every variant, with the poets scored by the same blind rubric. Human poets close on repair 19% of the time and on evasion 3%. OLMo 2 Instruct almost exactly inverts that at 4% and 69%.

Paired on the 87 poems both OLMo checkpoints finished admissibly, the direction of the change is specific (Table 2). Evasion rises 37.9 points and imposed hope 40.2 points, both at $p < 0.0001$ under a paired resampling test. Repair moves 1.1 points, $p = 1.00$. Repair and acceptance together move 0.0 points. The aligned model is not writing more consoling poems. It is writing poems that decline to stay with the loss.

That distinction is the reason the taxonomy exists. A valence score sees a brighter close and cannot say whether the brightness was earned from the poem's own materials or supplied from outside it.

4.2 The effect accumulates across the pipeline

Evasion climbs at every rung: 36% at base, 57% after supervised fine-tuning, 57% after DPO, 69% at the shipped instruct model. Imposed hope climbs 25%, 45%, 52%, 62%. Supervised fine-tuning on instruction data does most of the initial work, and preference optimization and reinforcement each

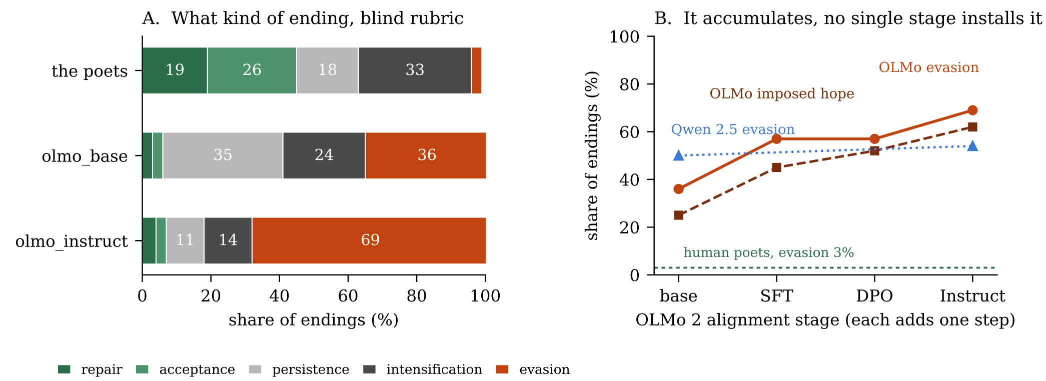

Figure 1: Left: the poets end with repair 19% of the time and evasion 3%. OLMo 2 Instruct reverses that to 4% and 69%, and its base checkpoint sits between them. Right: evasion and unearned hope climb at every rung of the OLMo ladder, while Qwen 2.5's instruction tuning moves neither. The key take-away is that alignment does not add consolation, it adds turning away.

Table 1: Share of endings by move, blind rubric, raw arm. The poets' endings were judged as unlabelled items alongside the models'. *Imposed hope* is scored independently of the move. The take-away is in the two boldface columns: every model rung has a lower repair rate than the poets and a far higher evasion rate.

variant	n	**repair**	acceptance	persistence	intensification	**evasion**	imposed hope
the poets	200	**19%**	26%	18%	33%	**3%**	4%
olmo_base	112	3%	3%	35%	24%	36%	25%
olmo_sft	58	2%	0%	21%	21%	57%	45%
olmo_dpo	81	6%	5%	14%	19%	57%	52%
olmo_instruct	155	**4%**	3%	11%	14%	**69%**	62%
qwen_base	132	3%	2%	27%	18%	50%	39%
qwen_instruct	140	3%	1%	16%	26%	54%	43%

add more. No single stage installs the disposition, which is the decomposition only a staged open release supports.

The valence metric agrees on ordering. On the 21 items every rung finished, mean ΔU against the poet is $+0.025$, $+0.108$, $+0.162$, $+0.207$ under VADER and $+0.291$, $+0.487$, $+0.551$, $+0.717$ under SST-2. Both orderings are monotonic and in the predicted direction, which happens with probability $1/4! = 0.042$ under a random ordering. The complete-case restriction matters here: scoring each rung on its own usable set makes the ladder appear to jump at DPO, and that jump is a selection artifact, since weaker rungs were being graded on the easier poems they managed to finish.

4.3 The claim is about a pipeline, not about alignment

Qwen 2.5's instruction tuning moves nothing. Evasion goes 54% to 54%, imposed hope 40% to 40%, repair 3% to 3%, every $p = 1.00$. The identical marginals look like a defect and are not: the two checkpoints assign different labels on 51 of the 99 shared items, and the distribution still lands in the same place. Tuning reshuffles which poems get which ending without shifting the population.

Both base models already evade heavily, at 36% and 50% against 3% for the poets, so pretraining supplies a large part of the disposition. What differs is what alignment does on top. OLMo 2 ships the full Tülu 3 recipe with documented stages. Qwen 2.5's instruction tuning is lighter and undocumented. Reporting this null is the reason the second family was run.

Table 2: Paired change from base to instruct within each family, raw arm, shared items only. The take-away is the contrast between the first two rows: what alignment adds is evasion and unearned hope, not repair, and Qwen's pipeline adds neither.

family	measure	base $\rightarrow$ instruct	change	p
OLMo 2 1B ($n = 87$)	evasion rate	$36\% \rightarrow 74\%$	$+37.9$ pp	< 0.0001
	imposed hope	$23\% \rightarrow 63\%$	$+40.2$ pp	< 0.0001
	repair rate	$0\% \rightarrow 1\%$	$+1.1$ pp	1.00
	repair or acceptance	$3\% \rightarrow 3\%$	$+0.0$ pp	1.00
	intensification (closes darker)	$23\% \rightarrow 10\%$	-12.6 pp	0.044
Qwen 2.5 1.5B ($n = 99$)	evasion rate	$54\% \rightarrow 54\%$	$+0.0$ pp	1.00
	imposed hope	$40\% \rightarrow 40\%$	$+0.0$ pp	1.00
	repair rate	$3\% \rightarrow 3\%$	$+0.0$ pp	1.00

4.4 Two routes to a brighter ending

Crossing the weight-level test against the interface test separates two mechanisms, and each family shows exactly one of them.

family	weights (base $\rightarrow$ instruct, raw)	interface (chat $-$ raw, same weights)
OLMo 2 1B	$+0.090$, $p = 0.018$ ($n = 87$)	$+0.026$, $p = 0.45$ ($n = 71$)
Qwen 2.5 1.5B	$+0.020$, $p = 0.58$ ($n = 99$)	$+0.089$, $p = 0.0032$ ($n = 125$)

OLMo's alignment moved the weights. The uplift appears under a plain-text prompt that gives the model no cue it is meant to be an assistant, and its own chat template adds nothing further. Qwen's alignment did not move the weights, and the assistant frame lifts its endings anyway. One disposition is durable and written into parameters. The other is summoned by being addressed as a helper. A study running only one arm would have concluded that whichever family it looked at was the general case.

4.5 Valence uplift, and where it is fragile

The primary valence test is significant for OLMo under VADER ($+0.090$, $p = 0.018$) and not under SST-2 ($+0.177$, $p = 0.060$), with the two scorers correlating at $r = 0.46$ item by item. Same direction, larger effect, significance crossing 0.05 under only one instrument. The Qwen null holds under both. Against the poet, OLMo Instruct closes $+0.142$ brighter ($p < 0.0001$, $d = 0.41$, $n = 155$) where the poets close flat at -0.021. Continuation length does not explain any of it: within-rung Spearman correlations between length and uplift stay at $|\rho| \leq 0.19$ with a minimum p of 0.087.

We lead with the taxonomy rather than with valence, because valence is the weaker instrument and because the taxonomy says the thing valence structurally cannot, which is that the brightness was bought by evasion.

5 Limitations

The judge is a language model deciding whether language models impose hope, and its labels have not yet been validated against human annotators. A blind rating task and a scorer-validation script are in the repository and the ratings are not collected. Until they are, the taxonomy results should be read as one instrument's consistent verdict rather than as a measured construct. Judge self-reported confidence averages 0.66 to 0.74, with 8% to 20% of judgments below 0.6.

Scale is 1B and 1.5B, forced by a 16 GB laptop where OLMo 2's 7B base checkpoint alone is 29 GB. This design shows that the effect exists and where in a pipeline it enters. It cannot show how it scales.

The human baseline is historical. Public domain means pre-1930 poets, and part of any gap may be period style rather than species. Contemporary theme-matched poems cannot be redistributed, which is why the public-domain set exists at all.

The SFT and DPO rungs ran on 100 items against 200 for the others, which is why complete-case analyses fall to $n = 21$. Each item was sampled once at temperature 0.8, so within-model sampling variance is not separated from between-model differences.

6 What the poems are for

The measurement here is quantitative and the object is not. Hardy wrote a poem whose entire argument is that no one is listening, and a 1B model with three stages of alignment on top answered it with a guardian angel. Horton ended a Civil War elegy on "Then sadly fallen and forever down!" and the same model closed it on rising again in Christ. Neither continuation is incompetent verse. Both refuse the poem's premise.

An artist working with these systems is working with a collaborator that has a standing objection to endings that do not resolve, and that objection is not in the interface where it could be turned off. In OLMo's case it is in the parameters. It survives a prompt that never mentions assistance, help, or safety. The base checkpoint is the only reason we can see it, and base checkpoints are the part of the release that most laboratories no longer publish.

References

[1] Team OLMo (2025) 2 OLMo 2 Furious. Allen Institute for AI. `https://arxiv.org/abs/2501.00656`

[2] Lambert, N., Morrison, J., Pyatkin, V., et al. (2025) Tülu 3: Pushing Frontiers in Open Language Model Post-Training. `https://arxiv.org/abs/2411.15124`

[3] Qwen Team (2025) Qwen2.5 Technical Report. `https://arxiv.org/abs/2412.15115`

[4] Tian, Y., Huang, T., Liu, M., Jiang, D., Spangher, A., Chen, M., May, J. & Peng, N. (2024) Are Large Language Models Capable of Generating Human-Level Narratives? In *Proceedings of EMNLP 2024*. `https://arxiv.org/abs/2407.13248`

[5] Hutto, C.J. & Gilbert, E. (2014) VADER: A parsimonious rule-based model for sentiment analysis of social media text. In *Proceedings of ICWSM 2014*.

[6] Socher, R., Perelygin, A., Wu, J., Chuang, J., Manning, C.D., Ng, A.Y. & Potts, C. (2013) Recursive deep models for semantic compositionality over a sentiment treebank. In *Proceedings of EMNLP 2013*, pp. 1631–1642.

[7] Smith, B.H. (1968) *Poetic Closure: A Study of How Poems End.* Chicago: University of Chicago Press.

[8] Rafailov, R., Sharma, A., Mitchell, E., Ermon, S., Manning, C.D. & Finn, C. (2023) Direct Preference Optimization: Your Language Model is Secretly a Reward Model. In *Advances in Neural Information Processing Systems 36*.

[9] Ouyang, L., Wu, J., Jiang, X., et al. (2022) Training language models to follow instructions with human feedback. In *Advances in Neural Information Processing Systems 35*.

[10] Bhyravajjula, S., Walsh, M., Preus, A. & Antoniak, M. (2025) so much depends / upon / a whitespace: Why Whitespace Matters for Poets and LLMs. In *Proceedings of EMNLP 2025*, pp. 35156–35173. `https://aclanthology.org/2025.emnlp-main.1783/`