
When Does a Fourier Neural Operator Fail? Spectral Diagnostics for Low-Rank Surrogates of Reaction–Diffusion

Hyunjun Yi*
Seoul International School
hjyi2027@gmail.com

David J. Zhang
Millburn High School
davidjzhang2011@gmail.com

Wenhao Lu
Millburn High School
wenhaolu2027@gmail.com

Jonathan Liu
Millburn High School
jonliu3000@gmail.com

Spursh Deshpande
Crescenta Valley High School
spurshdeshpande@gmail.com

Abstract

Neural surrogates for PDE simulators are usually judged by one-step error, which hides the two failures that matter in practice: error accumulation under autoregressive rollout and collapse on unseen parameter regimes. We study a compact CP-factorised Fourier neural operator (LiteFNO) against a parameter-matched CNN on the Gray–Scott benchmark from The Well, asking not only whether the surrogate fails but why. (i) LiteFNO uses 22% fewer FLOPs than the CNN, but both its one-step edge and its rollout advantage in a single seed reverse under three-seed averaging; rollout error is instead governed by the fate of the high-frequency spectrum, which the CNN over-damps and LiteFNO over-excites. (ii) Causally ablating the lowest-magnitude retained Fourier modes shows they contribute little at one step yet dominate mid-horizon rollout error. (iii) In a preliminary leave-one-regime-out study, failure on a held-out Gray–Scott regime correlates with the share of simulator spatial variance at low wavenumber (Spearman $\rho = 0.94$ over six regimes, single seed), a statistic requiring no trained model. We also report noise and 6-bit quantisation robustness, CPU/GPU latency, and a deferral pilot in which ensemble disagreement captures essentially all of the oracle-achievable gain, while showing that out-of-distribution error is spread across steps rather than concentrated in a discardable minority. These results argue for spectral diagnostics, computed from simulator output alone, as a cheap pre-training predictor of surrogate reliability.

1 Introduction

Machine-learned emulators promise to replace expensive simulators, but a surrogate is only useful if it is reliable on the trajectories and regimes a scientist actually queries. Two failure modes recur: *rollout drift*, where small one-step errors compound into unphysical states, and *regime collapse*, where a surrogate degrades sharply on a neighbouring region of parameter space. Both are well known Brandstetter et al. [2022], Lippe et al. [2023], McCabe et al. [2023], yet benchmarks typically report a one-step error blind to them.

We take a diagnostic, representation-level view: an FNO’s truncated spectral weights *are* its representation of the dynamics, and both failure modes can be traced to what that representation keeps

*This work was conducted as part of AIscent Labs (<https://aiscent-research.github.io/>).




Figure 1: **Headline result (3 seeds).** Left: 30-step rollout VRMSE. Right: ratio of predicted to true energy above wavenumber $k=8$. The CNN steadily loses high-frequency energy; LiteFNO over-excites it with large seed variance. Neither surrogate preserves the spectrum, and the seed-averaged rollout ranking reverses the single-seed result in Appendix Fig. 7.

and discards. Using the Gray–Scott reaction–diffusion system from *The Well* Ohana et al. [2024], re-simulated across its six qualitative regimes (Appendix Fig. 4), we compare a compact Fourier neural operator Li et al. [2021] whose spectral weights are CP-factorised (LiteFNO) against a parameter-matched CNN, and ask what in the learned operator, and what in the data, determines when each breaks. Whether a learned representation transfers to a new physical regime, and whether that is predictable *before* training, bears directly on when a trained surrogate can be trusted in place of an imperfect simulator.

Contributions.

1. **An honest efficiency/rollout comparison.** LiteFNO needs 22% fewer FLOPs than the CNN, but both its single-seed one-step and rollout advantages reverse under three-seed averaging (Sec. 3). We report both and show the rollout behaviour is explained by high-frequency spectral drift.
2. **Causal mode ablation.** Zeroing the lowest-magnitude retained Fourier modes of a trained LiteFNO changes one-step error far less than it changes mid-horizon rollout error (which nearly triples), isolating the weak tail of the spectrum as the carrier of rollout drift (Sec. 4).
3. **A pre-training predictor of regime collapse.** In a preliminary leave-one-regime-out study, the held-out/seen error ratio correlates with the fraction of ground-truth spatial variance at low wavenumber ($\rho = +0.94$, six regimes, single seed), a quantity requiring no trained model (Sec. 5).
4. **Deployment and reliability pilots.** Noise and quantisation robustness, batch-1 latency on CPU/GPU as a function of resolution, and a deferral/calibration pilot in which ensemble disagreement matches the oracle but out-of-distribution error cannot be deferred away, and ensembles are overconfident exactly where they are most certain (Sec. 6).

2 Setup

Simulator and data. Gray–Scott models two reacting species (u, v) with feed rate F and kill rate k ; varying (F, k) produces the spots, worms, maze, spirals, bubbles and gliders regimes (Appendix Fig. 4). We use the configuration of *The Well* Ohana et al. [2024], re-simulated with the reference finite-difference solver so that ground-truth energy spectra are available at full resolution. Models are trained on one-step prediction $x_t \mapsto x_{t+1}$ over both fields and evaluated with variance-normalised RMSE (VRMSE, the benchmark’s metric) at one step and over 30-step autoregressive rollouts.

Models. *LiteFNO* is a 2-D FNO with the standard $K \times K$ mode truncation ($K = 16$) whose complex spectral weight tensor is replaced by a rank- r CP decomposition, reducing parameters from $O(C^2 K^2)$ to $O(r(2C + 2K))$. The *CNN* baseline is a residual convolutional network matched to the FNO’s parameter budget. LiteFNO has 8 layers of width 64 with $K=16$ modes and fractional CP rank 0.02; both models are trained on 32^2 fields ($4 \times$ downsampled from *The Well*) with AdamW, learning rate 10^{-3} , StepLR (halved every 100 epochs), batch 64, 200 epochs, MSE loss and seeds 0/1/2 for Secs. 3–4. Our CNN reproduces the published CNN points of *The Well* (Appendix Fig. 6); the published FNO-S sits at $\sim 10^7$ parameters, motivating the compressed variant.




Figure 2: **Causal mode ablation (3 seeds).** Zeroing the lowest-energy fraction of LiteFNO’s Fourier modes at inference leaves one-step error nearly unchanged but sharply increases error in the 6–12 step window, while the late window (13–30) is comparatively insensitive.

3 Efficiency and rollout: a single seed is not enough

Compression. Sweeping CP rank on Gray–Scott (Appendix Fig. 5, left) shows test VRMSE flat from $r=8$ to $r=64$ (≈ 0.03 – 0.04) with a sharp degradation only at $r=128$, where optimisation becomes unstable at the FNO’s native parameter count. Our headline LiteFNO uses fractional CP rank 0.02 (full rank $R=77$; 68 components carry 90% of weight energy across seeds, Appendix Fig. 8a), so the factorisation is well utilised. It costs 1.69×10^8 FLOPs per forward pass against the CNN’s 2.19×10^8 (22% fewer), at $1.8 \times$ the parameter count (180k vs 108k). In a single seed LiteFNO is slightly more accurate at one step (0.0125 vs 0.0152 VRMSE; Appendix Fig. 5, right), but over three seeds the CNN is better (0.0105 vs 0.0132), so we do not claim a one-step accuracy advantage; no Fourier mode is dead (all normalised mode magnitudes exceed 0.7, Appendix Fig. 8b).

Rollout. In a single seed, LiteFNO accumulated error more slowly over 30 steps (0.44 vs 0.62 at step 30; Appendix Fig. 7); three-seed averaging (Fig. 1) reverses this, as at one step: the CNN is better at every horizon (0.54 vs 0.64), with LiteFNO’s seed spread roughly twice as large. We claim no rollout advantage for spectral factorisation; what is stable across seeds is the mechanism in the right panel.

4 Mechanism: truncated modes carry rollout error

Spectral drift. Fig. 1 (right) tracks predicted-to-true energy above $k=8$ along the rollout. The CNN decays monotonically to ≈ 0.7 , low-pass filtering the Turing pattern; LiteFNO overshoots to 1.3 (seed band to 1.8), feeding excess high-frequency energy back and amplifying it. One-step spectra match truth to $k \approx 20$ (Appendix Fig. 9a); the discrepancy is purely a rollout phenomenon.

Causal ablation. To test whether the weakly-weighted modes are causally responsible, we zero the lowest-magnitude fraction of the trained LiteFNO’s modes at inference (Fig. 2). Keeping 10% of modes raises one-step VRMSE from 0.013 to 0.062— $4.7 \times$ relative, yet small in absolute terms next to rollout: over the same ablation, error in steps 6–12 rises 0.28 to 0.72, sharpest between 75% and 50% of modes kept, while the late window (13–30), already dominated by accumulated error, moves only 0.47 to 0.59. The weakest retained modes—far from dead (Appendix Fig. 8b)—carry most of the mid-horizon signal, so mode-count decisions should be made on rollout, not one-step, metrics.

5 Regime generalisation may be predictable from the simulator alone

We train LiteFNO on five Gray–Scott regimes and test on the sixth, for all six folds (Fig. 3). The held-out penalty is wildly uneven: spots and gliders generalise for free (ratio ≤ 1), maze and worms lose little, but spirals and bubbles roughly double in error; the gap does not shrink from epoch 10 to 30 (no budget artefact), and for bubbles is still growing ($1.15 \rightarrow 1.60 \rightarrow 1.87$ —a lower bound).

The right panel explains the ordering: for each regime we compute, from ground truth only, the fraction of spatial variance in Fourier modes with $k < 8$. On the 32^2 grid a 16-mode FNO retains most of the spectrum, so $k=8$ is not the literal truncation boundary but the scale separating coherent structure from pattern-scale detail, chosen before training and untuned. Regimes dominated by large-scale structure (spirals, bubbles) are the ones the surrogate cannot extrapolate to (Spearman



Figure 3: **Leave-one-regime-out.** Left: one-step VRMSE on the five training regimes (green) vs the held-out regime (red). Middle: the held-out/seen ratio does not close with more epochs. Right: the ratio is predicted by the share of ground-truth spatial variance below Fourier mode 8 (Spearman $\rho = +0.94$).

$\rho = +0.94$): the model learns a local rule that transfers among fine-grained regimes, while large-scale coherent structures have no counterpart in the training spectrum. Appendix Fig. 10 gives the per-regime spectra.

Three caveats: this is one seed, four trajectories per regime, and a six-point rank correlation; it *inverts* our pre-registered hypothesis (maze and spots, with the most energy above the cutoff, were predicted hardest and were the two easiest), so the low-mode share is post hoc—a hypothesis generated by these data, not confirmed by them; and VRMSE normalises by per-regime variance, which co-varies with low-mode energy, so the correlation should be re-checked on unnormalised error. If it survives those tests, the predictor is *free*: one short simulation per candidate (F, k) , no training.

6 Deployment robustness and a deferral pilot

Robustness, precision, latency. Both models tolerate input noise to 20 dB SNR (LiteFNO slightly better at 0 dB: 1.15 vs 1.38) and 6-bit weight quantisation, collapsing at 4 bits (Appendix Figs. 12a–12b). Training with matched degradation augmentation halves error at 100% severity but costs 177% on clean inputs and does not transfer to a held-out corruption (Appendix Fig. 13). Parameter count does not predict latency ($\rho=0.07$; the smallest CP model is the slowest at batch 1, 31 ms), but FLOPs do once CP factors are fused into a dense spectral tensor; the fusing speed-up follows the closed-form batch decay but at $1.8\times$ instead of the predicted $12\times$ (Appendix Fig. 14). With grid side N , LiteFNO latency scales as $N^{0.6}-N^{1.1}$ vs the CNN’s $N^{1.3}-N^{1.9}$, favouring the FNO at high resolution; GPU throughput still favours the CNN $6.4\times/\approx 3\times$ at batch 1/16–64 (Appendix Fig. 15).

Deferral and calibration. If regime collapse is detectable at inference, unsafe steps can be handed back to the simulator. In a small pilot (one seed, 4 trajectories/regime, 118 held-out steps, two folds), we rank steps by disagreement of a 4-member ensemble against an oracle deferring the highest-error steps. On bubbles, the fold with a real $1.9\times$ held-out gap, disagreement captures 99–100% of the oracle-achievable gain at every deferral rate (random captures none); the maze fold (Appendix Fig. 11) shows the same tracking but had no gap to close. Less comfortably, on bubbles even 80% deferral leaves error at $1.05\times$ in-distribution—and so does the oracle: OOD error is spread across most steps, so deferral is a detector, not a fix. The same ensemble is overconfident (95% intervals cover 81%/73% in/out of distribution, ECE 0.066/0.131) and worst in its most confident bin: achieved error there is $4.6\text{--}5.2\times$ claimed (Appendix Fig. 16).

7 Related work, limitations, conclusion

Neural operators Li et al. [2021], Kovachki et al. [2023] and factorised variants Tran et al. [2023], Kossaifi et al. [2023] target the trade-off we study; our contribution is the diagnostic. Rollout instability has been attacked via pushforward training Brandstetter et al. [2022], spectral refinement Lippe et al. [2023], and spectral-bias analyses McCabe et al. [2023], Chattopadhyay and Hassanzadeh [2023]; we add a causal ablation tying drift to weakly-weighted retained modes. OOD generalisation of operators is studied in Subramanian et al. [2023], Li et al. [2024]; our low-mode share needs no trained model Ohana et al. [2024]. Deferral is selective prediction Geifman and El-Yaniv [2017].

Limitations and conclusion. All results are on one 2-D system at 32^2 ; only Secs. 3–4 use three seeds, the regime/deferral/calibration studies are single-seed pilots, and the Sec. 5 predictor is a six-

point post-hoc correlation with an open normalisation confound; forced temporal structure (Appendix Fig. 17) will need a temporal analogue of the cutoff. The message: the error that matters hides in the modes one-step evaluation ignores, and a one-line spectral statistic is a candidate predictor of regime coverage.

References

- Johannes Brandstetter, Daniel Worrall, and Max Welling. Message passing neural pde solvers. In *ICLR*, 2022.
- Ashesh Chattopadhyay and Pedram Hassanzadeh. Long-term instabilities of deep learning-based digital twins of the climate system: The cause and a solution. *arXiv:2304.07029*, 2023.
- Yonatan Geifman and Ran El-Yaniv. Selective classification for deep neural networks. In *NeurIPS*, 2017.
- Jean Kossaifi, Nikola Kovachki, Kamyar Azizzadenesheli, and Anima Anandkumar. Multi-grid tensorized fourier neural operator for high-resolution pdes. *arXiv:2310.00120*, 2023.
- Nikola Kovachki, Zongyi Li, Burigede Liu, Kamyar Azizzadenesheli, Kaushik Bhattacharya, Andrew Stuart, and Anima Anandkumar. Neural operator: Learning maps between function spaces. *JMLR*, 24, 2023.
- Zongyi Li, Nikola Kovachki, Kamyar Azizzadenesheli, Burigede Liu, Kaushik Bhattacharya, Andrew Stuart, and Anima Anandkumar. Fourier neural operator for parametric partial differential equations. In *ICLR*, 2021.
- Zongyi Li, Hongkai Zheng, Nikola Kovachki, David Jin, Haoxuan Chen, Burigede Liu, Kamyar Azizzadenesheli, and Anima Anandkumar. Physics-informed neural operator for learning partial differential equations. *ACM/IMS Journal of Data Science*, 1(3), 2024.
- Phillip Lippe, Bastiaan Veeling, Paris Perdikaris, Richard Turner, and Johannes Brandstetter. Pde-refiner: Achieving accurate long rollouts with neural pde solvers. In *NeurIPS*, 2023.
- Michael McCabe, Peter Harrington, Shashank Subramanian, and Jed Brown. Towards stability of autoregressive neural operators. *Transactions on Machine Learning Research*, 2023.
- Ruben Ohana, Michael McCabe, Lucas Meyer, Rudy Morel, Fruzsina J. Agocs, Miguel Beneitez, Marsha Berger, Blakesly Burkhart, Stuart B. Dalziel, Drummond B. Fielding, Daniel Fortunato, Jared A. Goldberg, Keiya Hirashima, Yan-Fei Jiang, Rich R. Kerswell, Suryanarayana Maddu, Jonah Miller, Payel Mukhopadhyay, Stefan S. Nixon, Jeff Shen, Romain Watteaux, Bruno Régaldou-Saint Blancard, François Rozet, Liam H. Parker, Miles Cranmer, and Shirley Ho. The well: a large-scale collection of diverse physics simulations for machine learning. In *NeurIPS Datasets and Benchmarks*, 2024.
- Shashank Subramanian, Peter Harrington, Kurt Keutzer, Wahid Bhimji, Dmitriy Morozov, Michael W. Mahoney, and Amir Gholami. Towards foundation models for scientific machine learning: Characterizing scaling and transfer behavior. In *NeurIPS*, 2023.
- Alasdair Tran, Alexander Mathews, Lexing Xie, and Cheng Soon Ong. Factorized fourier neural operators. In *ICLR*, 2023.

A The six Gray–Scott regimes




Figure 4: The six Gray–Scott regimes in *The Well*, re-simulated at full resolution with the (F, k) values shown. Regimes differ in their low-wavenumber variance share, a candidate predictor of surrogate generalisation (Sec. 5).

B Reproduction of *The Well* baselines and compression sweep

Figure 6 reproduces the published frontier; Fig. 5 gives the CP-rank sweep and FLOP comparison referenced in Sec. 3; Fig. 7 is the single-seed rollout discussed there.




Figure 5: Left: CP-rank sweep; accuracy is flat until $r=128$. Right: single-seed FLOP/accuracy comparison; the one-step ordering reverses over three seeds (Sec. 3).




Figure 6: Parameter/accuracy frontier reproduced from *The Well* logs. Our CNN overlaps the published CNN points; the published FNO-S variants lie near 10^7 parameters.




Figure 7: Single-seed rollout and spectral drift; compare with the 3-seed Fig. 1, where the rollout ordering reverses.

C Additional mechanism figures

Figures 8a–10 support Secs. 3–5: the trained LiteFNO uses its full CP rank and all Fourier modes, one-step spectra and error maps are nearly identical across models, and the per-regime spectra give the low-mode share used as the generalisation predictor.




Figure 8: Trained LiteFNO uses its CP rank and all Fourier modes.




Figure 9: One-step behaviour is nearly identical between models; differences are confined to pattern boundaries.




Figure 10: Per-regime spatial and temporal spectra and the variance retained below the mode cutoff used in Fig. 3.

D Robustness and deployment

Figures in this section support Sec. 6.





Figure 11: **Deferral pilot (held-out regime: maze)**. Left: retained VRMSE vs fraction of steps deferred; ensemble disagreement tracks the oracle to within 0.002 up to 60% deferral. Note the held-out error already lies below the in-distribution line at 0% deferral on this fold. Right: sensitivity/specificity of flagging unsafe steps.



(a) Input noise.

(b) Weight quantisation.

Figure 12: Both models tolerate 20 dB SNR and 6-bit weights.




Figure 13: Structured degradation: augmentation-matched training helps on the matched corruption, barely on the held-out one.




Figure 14: Latency study. (1) Parameters do not predict batch-1 latency; (2) FLOPs do once CP factors are fused; (3) fusing speed-up follows the closed-form batch decay but at a much smaller magnitude; (4) scaling with grid side N .




Figure 15: Inference throughput vs batch size.




Figure 16: Ensemble calibration, in-distribution and on held-out bubbles: overconfident at all nominal levels, and most overconfident in the highest-confidence bin.




Figure 17: Spatio-temporal variance budget of PlanetSWE; forced diurnal/annual lines dominate, motivating a temporal analogue of the spatial cutoff predictor.